

Tiny RISC-V Instruction Set Architecture

- Author : Christopher Batten
- Date : September 5, 2016

The Tiny RISC-V ISA is a subset of the 32-bit RISC-V ISA suitable for teaching. More specifically, the Tiny RISC-V ISA is a subset of the RV32IM ISA. The Tiny RISC-V ISA is divided into two versions: TinyRV1 includes just eight instructions and is suitable for lecture notes, problem sets, and exams, while TinyRV2 includes 34 instructions and is suitable for running simple C programs. This document provides a compact description of the Tiny RISC-V ISA, but it should be read in combination with the full RISC-V ISA manuals:

- <http://www.csl.cornell.edu/courses/ece4750/resources/waterman-riscv-isa-vol1-2p1.pdf>
- <http://www.csl.cornell.edu/courses/ece4750/resources/waterman-riscv-isa-vol2-1p9.pdf>

The full RISC-V ISA manuals include a wealth of useful information about why the architecture designers made specific design decisions.

Table of Contents

1. Architectural State
2. Tiny RISC-V Instruction Overview
3. Tiny RISC-V Instruction Encoding
 - 3.1. Instruction Formats
 - 3.2. Immediate Formats
4. Tiny RISC-V Instruction Details
 - 4.1. Control/Status Register Instructions
 - 4.2. Register-Register Arithmetic Instructions
 - 4.3. Register-Immediate Arithmetic Instructions
 - 4.4. Memory Instructions
 - 4.5. Unconditional Jump Instructions
 - 4.6. Conditional Branch Instructions
5. Tiny RISC-V Privileged ISA
6. Tiny RISC-V Pseudo-Instructions

1. Architectural State

Data Formats

- Tiny RISC-V only supports 4B signed and unsigned integer values. There are no byte nor half-word values and no floating-point.

General Purpose Registers

- There are 31 general-purpose registers x1-x31 (called x registers), which hold integer values. Register x0 is hardwired to the constant zero. Each register is 32 bits wide. Tiny RISC-V uses the same calling convention and symbolic register names as RISC-V:

+ x0	: zero	the constant value 0
+ x1	: ra	return address (caller saved)
+ x2	: sp	stack pointer (caller saved)
+ x3	: gp	global pointer
+ x4	: tp	thread pointer
+ x5	: t0	temporary registers (caller saved)
+ x6	: t1	"
+ x7	: t2	"
+ x8	: s0/fp	saved register or frame pointer (callee saved)
+ x9	: s1	saved register (callee saved)
+ x10	: a0	function arguments and/or return values (caller saved)

```

+ x11 : a1      "
+ x12 : a2      function arguments (caller saved)
+ x13 : a3      "
+ x14 : a4      "
+ x15 : a5      "
+ x16 : a6      "
+ x17 : a7      "
+ x18 : s2      saved registers (callee saved)
+ x19 : s3      "
+ x20 : s4      "
+ x21 : s5      "
+ x22 : s6      "
+ x23 : s7      "
+ x24 : s8      "
+ x25 : s9      "
+ x26 : s10     "
+ x27 : s11     "
+ x28 : t3      temporary registers (caller saved)
+ x29 : t4      "
+ x30 : t5      "
+ x31 : t6      "

```

Memory

- Tiny RISC-V only supports a 1MB virtual memory address space from 0x00000000 to 0x000fffff. The result of memory accesses to addresses larger than 0x000fffff are undefined.
- A key feature of any ISA is identifying the endianness of the memory system. Endianness specifies if we load a word in memory, what order should those bytes appear in the destination register. Assume the letter A is at byte address 0x0, the letter B is at byte address 0x1, the letter C is at byte address 0x2, and the letter D is at byte address 0x3. If we load a four-byte word from address 0x0, there are two options: the destination register can either hold 0xABCD (big endian) or 0xDCBA (little endian). There is no significant benefit of one system over the other. Tiny RISC-V uses a little endian memory system.

2. Tiny RISC-V Instruction Overview

Here is a brief list of the instructions which make up both versions of the Tiny RISC-V ISA, and then some discussion about the differences between the two versions.

TinyRV1

TinyRV1 contains a very small subset of the TinyRV2 ISA suitable for illustrating how small assembly sequences execute on various microarchitectures in lecture, problem sets, and exams.

- ADD, ADDI, MUL
- LW, SW
- JAL, JR
- BNE

TinyRV2

TinyRV2 is suitable for executing simple C programs that do not use system calls.

- ADD, ADDI, SUB, MUL
- AND, ANDI, OR, ORI, XOR, XORI
- SLT, SLTI, SLTU, SLTIU

- SRA, SRAI, SRL, SRLI, SLL, SLLI
- LUI, AUIPC
- LW, SW
- JAL, JALR
- BEQ, BNE, BLT, BGE, BLTU, BGEU
- CSRR, CSRW (proc2mgr, mgr2proc, stats_en, core_id, num_cores)

Discussion

TinyRV1 includes the JR instruction, but technically this is not a real instruction but is instead a pseudo-instruction for the following usage of JALR:

```
jalr x0, rs1, 0
```

The JALR instruction is a bit complicated, and we really only need the JR functionality to explain function calls in TinyRV1. So TinyRV1 only includes the JR pseudo-instruction, while TinyRV2 includes the full JALR instruction.

CSRR and CSRW are also pseudo-instructions in the full RV32IM ISA for specific usage of the CSRRW and CSRRS instructions. The full CSRRW and CSRRS instructions are rather complicated and we don't actually need any functionality beyond what CSRR and CSRW provide. So TinyRV2 only includes the CSRR and CSRW pseudo-instruction.

3. Tiny RISC-V Instruction Encoding

The Tiny RISC-V ISA uses the same instruction encoding as RISC-V. There are four instruction types and five immediate encodings. Each instruction has a specific instruction type, and if that instruction includes an immediate, then it will also have an immediate type.

3.1. Tiny RISC-V Instruction Formats

R-type

31	25 24	20 19	15 14	12 11	7 6	0	
+-----+-----+-----+-----+-----+-----+							
funct7		rs2		rs1		funct3 rd opcode	
+-----+-----+-----+-----+-----+-----+							

I-type

31	20 19	15 14	12 11	7 6	0
+-----+-----+-----+-----+-----+					
imm		rs1		funct3 rd opcode	
+-----+-----+-----+-----+-----+					

S-type

31	25 24	20 19	15 14	12 11	7 6	0	
+-----+-----+-----+-----+-----+							
imm		rs2		rs1		funct3 imm opcode	
+-----+-----+-----+-----+-----+							

U-type


31	11	7 6	0
+-----+-----+-----+			
imm		rd opcode	
+-----+-----+-----+			

3.2. Tiny RISC-V Immediate Formats


RISC-V has an asymmetric immediate encoding which means that the immediates are formed by concatenating different bits in an asymmetric order based on the specific immediate formats. Note that in RISC-V all immediates are always sign extended, and the sign-bit for the immediate is always in bit 31 of the instruction.

The following diagrams illustrate how to create a 32-bit immediate from each of the five immediate formats. The fields are labeled with the instruction bits used to construct their value. <-- n is used to indicate repeating bit n of the instruction to fill that field and z is used to indicate a bit which is always set to zero.


I-immediate




S-immediate




B-immediate



U-immediate



J-immediate



4. Tiny RISC-V Instruction Details

For each instruction we include a brief summary, assembly syntax, instruction semantics, instruction and immediate encoding format, and the actual encoding for the instruction. We use the following conventions when specifying the instruction semantics:

- R[rx] : general-purpose register value for register specifier rx
- CSR[src] : control/status register value for register specifier csr
- sext : sign extend to 32 bits
- M_4B[addr] : 4-byte memory value at address addr
- PC : current program counter
- <s : signed less-than comparison
- >=s : signed greater than or equal to comparison
- <u : unsigned less-than comparison
- >=u : unsigned greater than or equal to comparison

- imm : immediate according to the immediate type

Unless otherwise specified assume instruction updates PC with PC+4.

4.1. Control/Status Register Instructions

CSRR

- Summary : Move value in control/status register to GPR
- Assembly : csrr rd, csr
- Semantics : $R[rd] = CSR[csr]$
- Format : I-type, I-immediate

31	20 19	15 14	12 11	7 6	0
-----+	-----+	-----+	-----+	-----+	-----+
csr	rs1	010	rd	1110011	
-----+	-----+	-----+	-----+	-----+	-----+

The control/status register read instruction is used to read a CSR and write the result to a GPR. The CSRs supported in TinyRV2 are listed in Section 5. Note that in RISC-V CSRR is really a pseudo-instruction for a specific usage of CSRRS, but in TinyRV2 we only support the subset of CSRRS captured by CSRR. See Section 6 for more details about pseudo-instructions.

CSRW

- Summary : Move value in GPR to control/status register
- Assembly : csrw csr, rs1
- Semantics : $CSR[csr] = R[rs1]$
- Format : I-type, I-immediate

31	20 19	15 14	12 11	7 6	0
-----+	-----+	-----+	-----+	-----+	-----+
csr	rs1	001	rd	1110011	
-----+	-----+	-----+	-----+	-----+	-----+

The control/status register write instruction is used to read a GPR and write the result to a CSR. The CSRs supported in TinyRV2 are listed in Section 5. Note that in RISC-V CSRW is really a pseudo-instruction for a specific usage of CSRRW, but in TinyRV2 we only support the subset of CSRRW captured by CSRW. See Section 6 for more details about pseudo-instructions.

4.2. Register-Register Arithmetic Instructions

ADD

- Summary : Addition with 3 GPRs, no overflow exception
- Assembly : add rd, rs1, rs2
- Semantics : $R[rd] = R[rs1] + R[rs2]$
- Format : R-type

31	25 24	20 19	15 14	12 11	7 6	0
-----+	-----+	-----+	-----+	-----+	-----+	-----+
0000000	rs2	rs1	000	rd	0110011	
-----+	-----+	-----+	-----+	-----+	-----+	-----+

SUB

- Summary : Subtraction with 3 GPRs, no overflow exception
- Assembly : sub rd, rs1, rs2
- Semantics : $R[rd] = R[rs1] - R[rs2]$
- Format : R-type

31	25 24	20 19	15 14	12 11	7 6	0
0100000	rs2	rs1	000	rd	0110011	

AND

- Summary : Bitwise logical AND with 3 GPRs
- Assembly : and rd, rs1, rs2
- Semantics : $R[rd] = R[rs1] \& R[rs2]$
- Format : R-type

31	25 24	20 19	15 14	12 11	7 6	0
0000000	rs2	rs1	111	rd	0110011	

OR

- Summary : Bitwise logical OR with 3 GPRs
- Assembly : or rd, rs1, rs2
- Semantics : $R[rd] = R[rs1] | R[rs2]$
- Format : R-type

31	25 24	20 19	15 14	12 11	7 6	0
0000000	rs2	rs1	110	rd	0110011	

XOR

- Summary : Bitwise logical XOR with 3 GPRs
- Assembly : xor rd, rs1, rs2
- Semantics : $R[rd] = R[rs1] \wedge R[rs2]$
- Format : R-type

31	25 24	20 19	15 14	12 11	7 6	0
0000000	rs2	rs1	100	rd	0110011	

SLT

- Summary : Record result of signed less-than comparison with 2 GPRs
- Assembly : slt rd, rs1, rs2
- Semantics : $R[rd] = (R[rs1] <_s R[rs2])$
- Format : R-type

31	25 24	20 19	15 14	12 11	7 6	0
0000000	rs2	rs1	010	rd	0110011	

This instruction uses a `_signed_` comparison.

SLTU

- Summary : Record result of unsigned less-than comparison with 2 GPRs
- Assembly : sltu rd, rs1, rs2
- Semantics : $R[rd] = (R[rs1] <_u R[rs2])$
- Format : R-type


31	25 24	20 19	15 14	12 11	7 6	0
0000000	rs2	rs1	011	rd	0110011	

```
+-----+-----+-----+-----+-----+-----+
```

This instruction uses an `_unsigned_` comparison.

SRA


- Summary : Shift right arithmetic by register value (sign-extend)
- Assembly : `sra rd, rs1, rs2`
- Semantics : $R[rd] = R[rs1] \ggg R[rs2][4:0]$
- Format : R-type



Note that the hardware should ensure that the sign-bit of `R[rs1]` is extended to the right as it does the right shift. The hardware `_must_` only use the bottom five bits of `R[rs2]` when performing the shift.

SRL


- Summary : Shift right logical by register value (append zeroes)
- Assembly : `srl rd, rs1, rs2`
- Semantics : $R[rd] = R[rs1] \gg R[rs2][4:0]$
- Format : R-type



Note that the hardware should append zeros to the left as it does the right shift. The hardware `_must_` only use the bottom five bits of `R[rs2]` when performing the shift.

SLL


- Summary : Shift left logical by register value (append zeroes)
- Assembly : `sll rd, rs1, rs2`
- Semantics : $R[rd] = R[rs1] \ll R[rs2][4:0]$
- Format : R-type



Note that the hardware should append zeros to the right as it does the left shift. The hardware `_must_` only use the bottom five bits of `R[rs2]` when performing the shift.

MUL


- Summary : Signed multiplication with 3 GPRs, no overflow exception
- Assembly : `mul rd, rs1, rs2`
- Semantics : $R[rd] = R[rs1] * R[rs2]$
- Format : R-type



4.3. Register-Immediate Arithmetic Instructions


ADDI

- Summary : Add constant, no overflow exception
- Assembly : addi rd, rs1, imm
- Semantics : $R[rd] = R[rs1] + sext(imm)$
- Format : I-type, I-immediate




ANDI

- Summary : Bitwise logical AND with constant
- Assembly : andi rd, rs1, imm
- Semantics : $R[rd] = R[rs1] \& sext(imm)$
- Format : I-type, I-immediate




ORI

- Summary : Bitwise logical OR with constant
- Assembly : ori rd, rs1, imm
- Semantics : $R[rd] = R[rs1] | sext(imm)$
- Format : I-type, I-immediate




XORI

- Summary : Bitwise logical XOR with constant
- Assembly : xori rd, rs1, imm
- Semantics : $R[rd] = R[rs1] \wedge sext(imm)$
- Format : I-type, I-immediate



SLTI


- Summary : Set GPR if source GPR < constant, signed comparison
- Assembly : slti rd, rs1, imm
- Semantics : $R[rd] = (R[rs1] <_s sext(imm))$
- Format : I-type, I-immediate



SLTIU


- Summary : Set GPR if source GPR is < constant, unsigned comparison
- Assembly : sltiu rd, rs1, imm

- Semantics : $R[rd] = (R[rs1] \ll sext(imm))$
- Format : I-type, I-immediate



SRAI


- Summary : Shift right arithmetic by constant (sign-extend)
- Assembly : `srai rd, rs1, imm`
- Semantics : $R[rd] = R[rs1] \ggg imm$
- Format : I-type



Note that the hardware should ensure that the sign-bit of $R[rs1]$ is extended to the right as it does the right shift.

SRLI


- Summary : Shift right logical by constant (append zeroes)
- Assembly : `srlr rd, rs1, imm`
- Semantics : $R[rd] = R[rs1] \gg imm$
- Format : I-type



Note that the hardware should append zeros to the left as it does the right shift.

SLLI


- Summary : Shift left logical constant (append zeroes)
- Assembly : `slli rd, rs1, imm`
- Semantics : $R[rd] = R[rs1] \ll imm$
- Format : R-type



Note that the hardware should append zeros to the right as it does the left shift.


LUI

- Summary : Load constant into upper bits of word
- Assembly : `lui rd, imm`
- Semantics : $R[rd] = imm \ll 12$
- Format : I-type, U-immediate



AUIPC


- Summary : Load PC + constant into upper bits of word
- Assembly : auipc rd, imm
- Semantics : $R[rd] = PC + (imm \ll 12)$
- Format : I-type, U-immediate



4.4. Memory Instructions

LW


- Summary : Load word from memory
- Assembly : lw rd, imm(rs1)
- Semantics : $R[rd] = M_4B[R[rs1] + sext(imm)]$
- Format : I-type, I-immediate



All addresses used with LW instructions must be four-byte aligned. This means the bottom two bits of every effective address (i.e., after the base address is added to the offset) will always be zero.

SW

- Summary : Store word into memory
- Assembly : sw rs2, imm(rs1)
- Semantics : $M_4B[R[rs1] + sext(imm)] = R[rs2]$
- Format : S-type, S-immediate




All addresses used with SW instructions must be four-byte aligned. This means the bottom two bits of every effective address (i.e., after the base address is added to the offset) will always be zero.

4.5. Unconditional Jump Instructions

JAL

- Summary : Jump to address and place return address in GPR
- Assembly : jal rd, imm
- Semantics : $R[rd] = PC + 4; PC = PC + sext(imm)$
- Format : U-type, J-immediate



JR

- Summary : Jump to address
- Assembly : jr rs1
- Semantics : $PC = R[rs1]$

- Format : I-Type, I-immediate

31	20 19	15 14	12 11	7 6	0
000000000000	rs1	000	00000	1100111	

Note that JR is a "real" instruction in TinyRV1, but it is a pseudo-instruction for a specific usage of JALR. We don't really worry about zero-ing out the the least-significant bit to zero in TinyRV1, but this must be done for TinyRV2.

JALR

- Summary : Jump to address and place return address in GPR
- Assembly : jalr rd, rs1, imm
- Semantics : $R[rd] = PC + 4; PC = (R[rs1] + sext(imm)) \& 0xfffffffffe$
- Format : I-Type, I-immediate

31	20 19	15 14	12 11	7 6	0
imm	rs1	000	rd	1100111	

Note that the target address is obtained by adding the 12-bit signed I-immediate to the value in register rs1, then setting the least-significant bit of the result to zero. In other words, the JALR instruction ignores the lowest bit of the calculated target address.

4.6. Conditional Branch Instructions

BEQ

- Summary : Branch if 2 GPRs are equal
- Assembly : beq rs1, rs2, imm
- Semantics : $PC = (R[rs1] == R[rs2]) ? PC + sext(imm) : PC + 4$
- Format : S-type, B-immediate

31	25 24	20 19	15 14	12 11	7 6	0
imm	rs2	rs1	000	imm	1100011	

BNE

- Summary : Branch if 2 GPRs are not equal
- Assembly : bne rs1, rs2, imm
- Semantics : $PC = (R[rs1] != R[rs2]) ? PC + sext(imm) : PC + 4$
- Format : S-type, B-immediate

31	25 24	20 19	15 14	12 11	7 6	0
imm	rs2	rs1	001	imm	1100011	

BLT

- Summary : Branch based on signed comparison of two GPRs
- Assembly : blt rs1, rs2, imm
- Semantics : $PC = (R[rs1] <_s R[rs2]) ? PC + sext(imm) : PC + 4$
- Format : S-type, B-immediate


31	25 24	20 19	15 14	12 11	7 6	0
imm	rs2	rs1	100	imm	1100011	

+-----+-----+-----+-----+-----+-----+

This instruction uses a `_signed_` comparison.

BGE


- Summary : Branch based on signed comparison of two GPRs
- Assembly : `bge rs1, rs2, imm`
- Semantics : $PC = (R[rs1] \geq_s R[rs2]) ? PC + sext(imm) : PC + 4$
- Format : S-type, B-immediate



This instruction uses a `_signed_` comparison.

BLTU


- Summary : Branch based on unsigned comparison of two GPRs
- Assembly : `bltu rs1, rs2, imm`
- Semantics : $PC = (R[rs1] <_u R[rs2]) ? PC + sext(imm) : PC + 4$
- Format : S-type, B-immediate



This instruction uses an `_unsigned_` comparison.

BGEU

- Summary : Branch based on unsigned comparison of two GPRs
- Assembly : `bgeu rs1, rs2, imm`
- Semantics : $PC = (R[rs1] \geq_u R[rs2]) ? PC + sext(imm) : PC + 4$
- Format : S-type, B-immediate



This instruction uses an `_unsigned_` comparison.

5. Tiny RISC-V Privileged ISA

Tiny RISC-V does not support any kind of distinction between user and privileged mode. Using the terminology in the RISC-V vol 2 ISA manual, Tiny RISC-V only supports M-mode.

Reset Vector

- RISC-V specifies two potential reset vectors: one at a low address, and one at a high address. Tiny RISC-V uses the low address reset vector at `0x00000200`. This is where assembly tests should reside as well as user code in TinyRV2.

Control/Status Registers

RISC-V includes four non-standard CSRs. `numcores` has the same meaning as the standard CSR `mhartid`, so we make `numcores` a synonym for `mhartid`. Here is the mapping:

CSR Name	Privilege	Read/Write	CSR Num	Note
proc2mngr	M	RW	0x7C0	non-standard
mngr2proc	M	R	0xFC0	non-standard
coreid	M	R	0xF14	synonym of mhartid (hardware thread ID)
numcores	M	R	0xFC1	non-standard
stats_en	M	RW	0x7C1	non-standard

These are chosen to conform to the guidelines in Section 2.1 of the RISC-V vol 2 ISA manual. Here is a description of each of these five CSRs.

- mngr2proc: 0xFC0

Used to communicate data from the manager to the processor. This register has register-mapped FIFO-dequeue semantics meaning reading the register essentially dequeues the data from the head of a FIFO. Reading the register will stall if the FIFO has no valid data. Writing the register is undefined.

- proc2mngr: 0x7c0

Used to communicate data from the processor to the manager. This register has register-mapped FIFO-enqueue semantics meaning writing the register essentially enqueues the data on the tail of a FIFO. Writing the register will stall if the FIFO is not ready. Reading the register is undefined.

- stats_en: 0x7c1

Used to enable or disable the statistics tracking feature of the processor (i.e., counting cycles and instructions)

- numcores: 0xfc1

Used to store the number of cores present in a multi-core system. Writing the register is undefined.

- coreid: 0xf14

Used to communicate the core id in a multi-core system. Writing the register is undefined.

Address Translation

Tiny RISC-V only supports the most basic form of address translation. Every logical address is directly mapped to the corresponding physical address. As mentioned above, Tiny RISC-V only supports a 1MB virtual memory address space from 0x00000000 to 0x000fffff, and thus Tiny RISC-V only supports a 1MB physical memory address space. In the RISC-V vol 2 ISA manual this is called a Mbare addressing environment.

6. Tiny RISC-V Pseudo-Instructions

It is very important to understand the relationship between the "real" instructions presented in this manual, the "real" instructions in the official RISC-V ISA manual, and pseudo-instructions. There are four instructions we need to be careful with: NOP, JR, CSRR, CSRW. The following table illustrates which ISAs contain which of these four instructions, and whether or not the instruction is considered a "real" instruction or a "pseudo-instruction".

```

-----
NOP   pseudo  pseudo  pseudo
JR    real    pseudo  pseudo
CSRR  --      real    pseudo
CSRW  --      real    pseudo

```

NOP is always a pseudo-instruction. It is always equivalent to the following use of the ADDI instruction:

```
nop == addi x0, x0, 0
```

JR is a "real" instruction in TinyRV1, but it is a pseudo-instruction in TinyRV2 and RISC-V for the following use of the JALR instruction:

```
jr rs1 == jalr x0, rs1, 0
```

CSRR and CSSRW are real instructions in TinyRV2 but they are pseudo-instructions for the following use of the CSRRS and CSRRW:

```
csrr rd, csr == csrrs rd, csr, x0
csrw csr, rs1 == csrrw x0, csr, rs1
```

None of this changes the encodings. In TinyRV1, JR is encoded the same way as the corresponding use of the JALR instruction in TinyRV2.